



Lustre 2.16 and Beyond

Andreas Dilger

Lai Siyao

Trends in High Performance Storage

- ▶ Ever increasing demand for high-performance storage to feed data pipelines
- ▶ AI/ML/ChatGPT/LLM driving surge in new users of parallel (flash) storage
- ▶ Existing computation models (weather, finance, ...) increasing resolution, data sources, historical data
- ▶ Initial demands can be met by all-flash storage, but not everyone has the budget to scale flash
- ▶ Need to increase capacity, reduce costs, transparently access multiple storage types/tiers/hybrid
- ▶ Lustre already allows transparent data migration between tiers, hybrid storage *within* files
- ▶ Disk and QLC, combined with compression, heading to lowest \$/TB and more accessible than tape
- ▶ Meta/data redundancy improves availability above hardware, simplifies hardware requirements
- ▶ Security, multi-tenancy, data isolation demands always increasing (medical, privacy, IP, legislative, ...)

Planned Feature Release Highlights

▶ 2.16 approaching feature completion

- **LNet IPv6 addressing** – must-have functionality for future deployments (SuSE, ORNL)
- **Optimized Directory Traversal (WBC1)** – improve efficiency for accessing many files (WC)

▶ 2.17 has major features already well underway

- **Client-side data compression** – reduce network and storage usage, costs (WC, UHamburg)
- **Metadata Writeback Cache (WBC2)** – order of magnitude better metadata speed (WC)

▶ 2.18 feature proposals in early stages

- **File Level Redundancy - Erasure Coding (FLR-EC)** – reduce cost, improve availability (ORNL)
- **Lustre Metadata Redundancy (LMR1)** – improve availability for large DNE systems
- **Client Container Image (CCI)** – improved handling of aggregations of many small files

LNet Improvements

Demand for IPv6 in new deployments as IPv4 is exhausted

- Relatively few external-facing Lustre systems means 10.x.y.z is still viable for now

► IPv6 large NID support ([LU-10391](#) SuSE, ORNL)

- Variable-sized NIDs (8-bit LND type, 8-bit address size, 16-bit network number, 16-byte+ address)
- Interoperable with existing current LNDs whenever possible
- Enhancements to LNet/socklnd for large NIDs mostly finished
- Work ongoing to handle large NIDs in Lustre code
 - Mount, config logs, [Imperative Recovery](#), [Nodemaps](#), root squash, etc.

► Improved network discovery/peer health (HPE, WC)

► Simplified/dynamic server node addressing ([LU-14668](#) WC)


- Detect added/changed server interfaces automatically ([LU-10360](#))
- Reduce (eventually eliminate) static NIDs in Lustre config logs
- Simplified handling for IPv6 NIDs by clients



Client-Side Usability and Performance Improvements

Ongoing ease-of-use and performance improvements for users and admins

- ▶ Parallel file/directory rename within a directory ([LU-12125](#) WC)
- ▶ llstat, llobdstat easier to use ([LU-13705](#) WC)
- 2.15 ▶ lfs find -printf formatted output of specific fields ([LU-10378](#) ORNL)
- 2.16 ▶ lfs migrate performance improvements ([LU-16587](#) HPE, WC)
- ▶ lfs migrate bandwidth limit, progress updates ([LU-13482](#) Amazon)
- ▶ Ongoing code updates/cleanup for newer kernels (ORNL, HPE, SuSE)
- 2.17 ▶ Client-side Data Compression ([LU-10026](#) WC, UHamburg, Intel)
- ▶ Buffered/DIO performance/efficiency improvements ([LU-13805](#), [LU-14950](#) WC)
- ▶ Erasure Coded FLR files ([LU-10911](#) ORNL)



Client-Side Data Compression

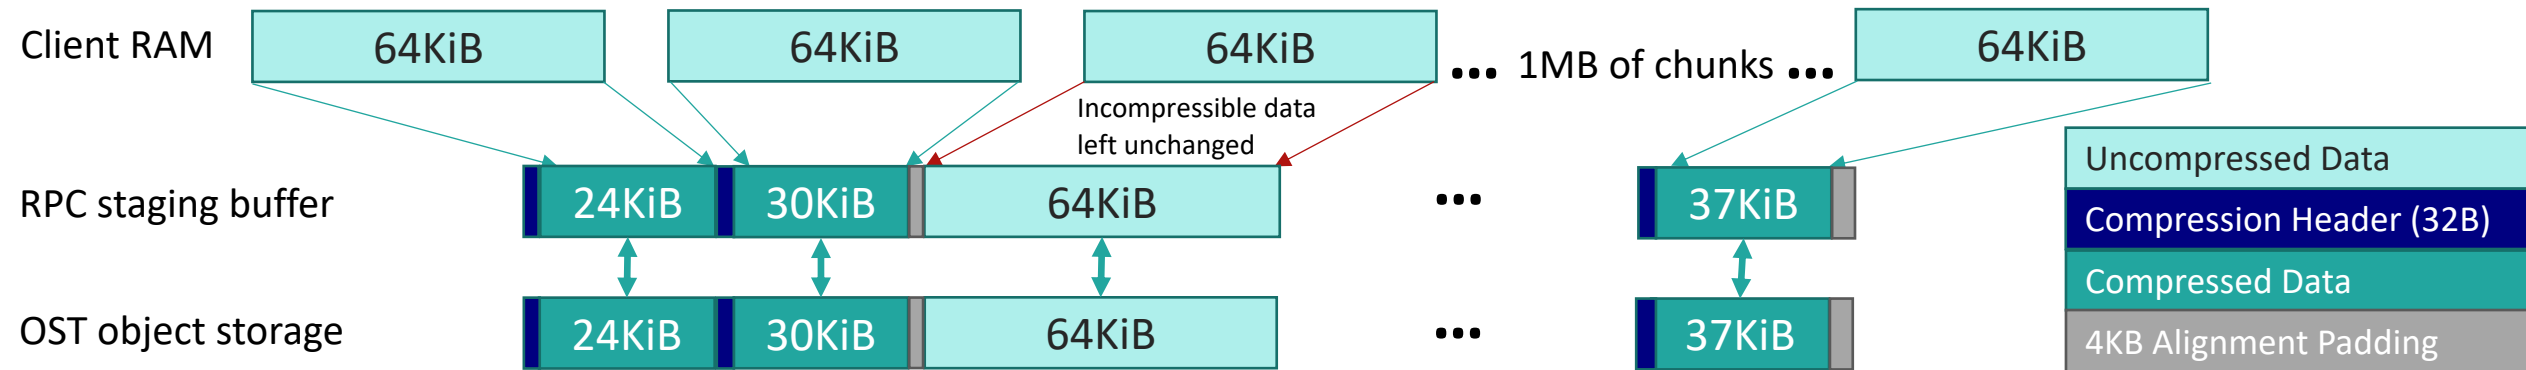
(WC, UHamburg 2.17+) 
Whamcloud

Increased capacity and lower cost for all-flash OSTs

- Parallel compression of RPCs on client cores for GB/s speeds, **no server CPU overhead!**

► (De-)Compress (lzo, lz4, gzip,...) RPC on client in chunks (64KiB-1MiB+) ([LU-10026](#))

- **Per directory or file component** selection of algorithm, level, chunk size (PFL, FLR)
- Keep "uncompressed" chunks as-is for incompressible data/file (.gz, .jpg, .mpg, ...)



- Client writes/reads whole chunk(s), (de-)compresses to/from RPC staging buffer
 - Larger chunks improve compression, but higher decompress/read-modify-write overhead
- Optional write uncompressed to one FLR mirror for random IO pattern
- Optional data (re-)compression during mirror/migrate to slow tier (via data mover)

Server-side Capacity and Efficiency Improvements



Ongoing performance and capacity scaling for next-gen hardware and systems

- ▶ OST object directory scalability for multi-PB OSTs ([LU-11912](#) WC)
 - Regularly create new object subdirectories (every 32M creates vs. 4B creates)
 - Better handling for billions of objects, grouping by age optimizes RAM and IOPS
- ▶ Read-only mounting of OST and MDT devices ([LU-15873](#) WC)
- ▶ Improved e2fsck for large dir and shared block errors ([LU-14710](#), [LU-16171](#) WC)
- ▶ lljobstat utility for easily monitoring "top" jobs ([LU-16228](#) WC)
 - 2.16 • Add IO size histograms to job_stats output, handle bad job names better
- 2.17 ▶ Reduced transaction size for many-striped files/dirs ([LU-14918](#) WC)
- ▶ Improved ldiskfs mballoc efficiency for large filesystems ([LU-14438](#) Google, WC, HPE)
- ▶ Parallel e2fsck for pass2/3 (directory entries, name linkage) ([LU-14679](#) WC)

Improved Data Security and Containerization

Growing dataset sizes and varied uses increases need to isolate users and their data

- ▶ Filenames encrypted on client in directory entries ([LU-13717](#) WC)
- ▶ Migrate/mirror of encrypted files without key ([LU-14667](#) WC)
- ▶ Encrypted file backup/restore/HSM without key ([LU-16374](#) WC)
- 2.15 ▶ Nodemap project quota mapping, squash all files to project ([LU-14797](#) WC)
- 2.16 ▶ Read-only mount enforced for nodemap clients ([LU-15451](#) WC)
- ▶ Kerberos authentication improvements ([LU-16630](#), [LU-16646](#) WC, NVIDIA)
- ▶ Nodemap Role-Based Admin Controls (fscrypt, changelog, chown, quota) ([LU-16524](#) WC)
- ▶ Cgroup/memcg memory usage limits for containers/jobs on clients ([LU-16671](#) WC, HPE)



Metadata Scaling Improvements

(WC 2.15+)



Improve usability and ease of DNE metadata horizontal performance/capacity scaling

► **DNE MDT Space Balance** - load balancing with normal mkdir ([LU-13417](#), [LU-13440](#))

- Round-robin/balanced subdirs, limited layout inheritance depth, less need for striped dirs

► Single-dir migration without recursion - "lfs migrate -m -d <dir>" ([LU-14975](#))

2.15 ► Balanced migration prefers less full MDTs - "lfs migrate -m -1 <dir>" ([LU-13076](#))

2.16 ► DNE inode migration improvements ([LU-14719](#), [LU-15720](#))

- Pre-check target space, stop on error, improved CRUSH2 hash

► More robust DNE MDT llog recovery ([LU-16203](#), [LU-16159](#))


- Handle errors and inconsistencies in recovery logs better

► DNE locking, remote RPC optimization ([LU-15528](#))

- Distributed transaction performance, reduce lock contention

2.17 ► Lustre Metadata Robustness/Redundancy ([LU-12310](#))

- Phase 1 to distribute/mirror MDT0000 services to other MDTs



Batched Cross-Directory Statahead (WBC1)

Improved access speed and efficiency for large directories/trees

- IO500 mdtest - {easy/hard} - stat performance improved 77%/95%

► **Batched RPC** infrastructure for multi-update operations ([LU-13045](#))


- Allow multiple getattrs/updates packed into a single MDS RPC
- More efficient network and server-side request handling

► **Batched statahead** for `ls -l`, `find`, etc. ([LU-14139](#))

- Aggregate getattr RPCs for existing statahead mechanism

► **Cross-Directory statahead** pattern matching ([LU-14380](#))

- Detect breadth-first (**BFS**) depth-first (**DFS**) directory tree walk
- Direct statahead to next file/subdirectory based on tree walk pattern
- Detect strided pattern for alphanumeric ordered traversal + `stat()`
 - e.g. `file00001, file001001, file002001...` or `file1, file17, file31, ...` order



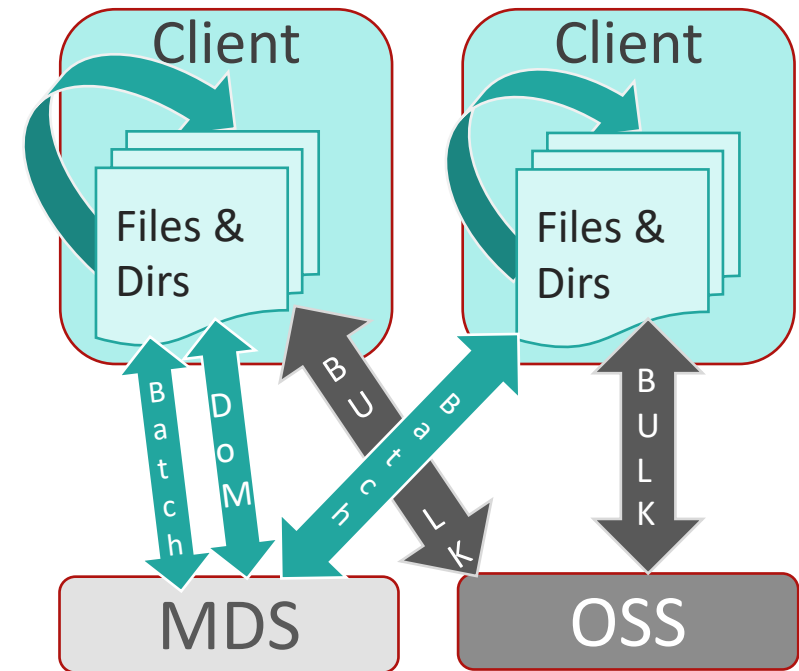
Metadata Writeback Cache (WBC2)

(WC 2.17+)



10-100x speedup for single-client file/dir create-intensive workloads

- Genome extraction/processing, untar/build, data ingest, producer/consumer
-
- ▶ Create new dirs/files **in client RAM without RPCs** ([LU-10983](#))
 - Lock **new** directory exclusively at mkdir time
 - Cache new files/dirs/data in RAM until cache flush or remote access
 - ▶ **No RPC round-trips** for file modifications in new directory
 - ▶ Batch RPC for efficient directory fetch and cache flush
 - ▶ **Files globally visible** on remote client access
 - Flush top-level entries, exclusively lock new subdirs, unlock parent
 - Repeat as needed for subdirectories being accessed remotely
 - Flush rest of tree in background to MDS/OSS by age or size limits
 - ▶ Productization of WBC code well underway
 - Some complexity handling partially-cached directories
 - Able to benchmark under intensive multi-client workloads
 - ▶ WBC for pre-existing directories, PCC integration in later release
 - Read all directory entries before create to avoid duplicate filenames



Enhance Dir Migration ([LU-17148](#))

(2.17)



Enhance directory migration robustness

- ▶ Directory migration relies on recovery to avoid file missing upon failure
- ▶ Once recovery is aborted, a different mechanism should be applied:
 - Source inodes and dirents of sub files are kept till migration of this directory finishes
 - Upon a directory migration finish, sync directory and verify sub file targets against retained sources
 - If some files are missing, migrate them, and redo the above step
 - Otherwise remove retained sources
- ▶ MDT resumes failed migrations automatically

Lustre Metadata Redundancy ([LU-12310](#))

(2.17+)



Improve metadata (data) availability in face of network/server errors

- In early discussion and planning stages

► LMR1a: Redundant services on other MDTs

- Mirror FLDB, Quota, flock() across MDTs

► LMR1b: DNE transaction performance

- Need to optimize if all mkdir are mirrored transactions
- Better distributed transaction logging format
- Improves **all** DNE operation performance

2.17

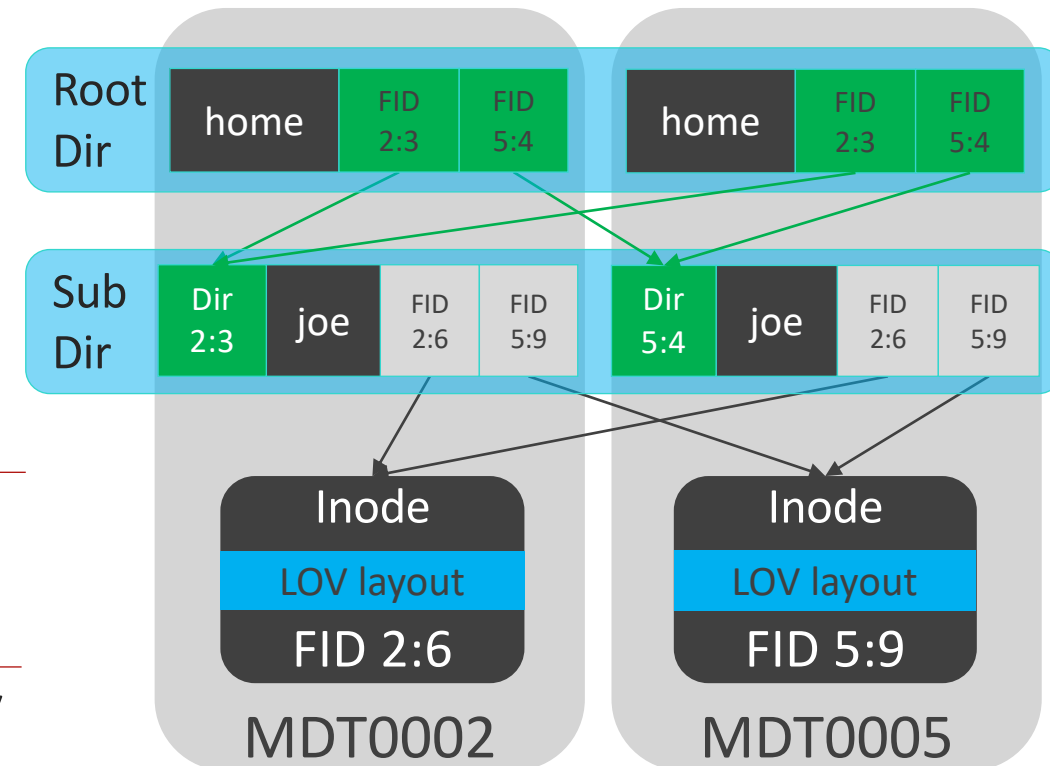
► LMR1c: Replicate top-level dirs for availability

- ROOT/ directory (rarely changed) mirrored to 2+ MDTs
- No replication for regular file inodes in this phase


2.18

2.19+ ► LMR2/3 phases needed for full MDT redundancy

- Full directory tree replication, file inode replication
- Configurable mirror setting per directory/tree (lfs setdirstripe)
- Recovery, LFSSCK, rebuild replicated directories if full MDT loss



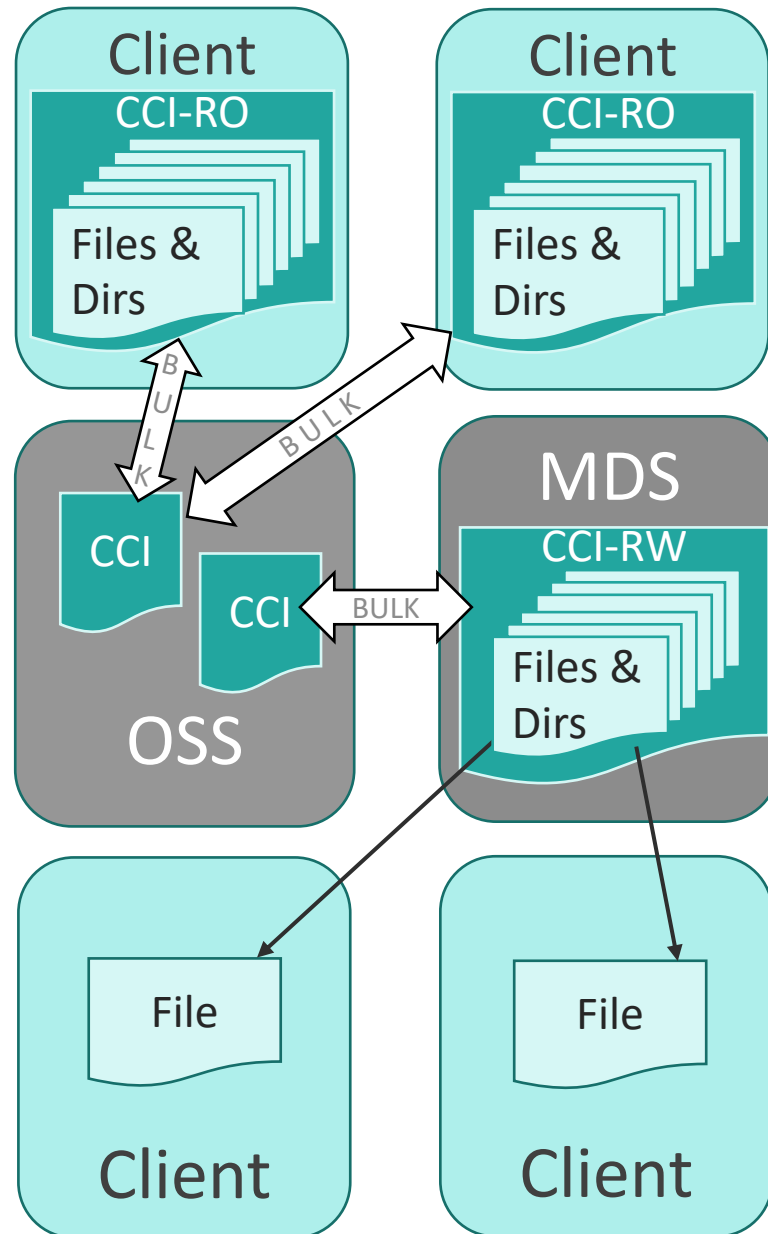
Client Container Image (CCI) (2.18+)



Need improved handling of aggregations of many files

- Create, access, (modify?), delete many files without untar/unzip
- ▶ CCI allows efficiently accessing filesystem image many files
- ▶ **Low I/O overhead, few file lock(s), high IOPS per client**
 - Readahead and write merging for data and metadata
 - Client-local in-RAM filesystem operations with very low latency
- ▶ Access, migrate, replicate image with large bulk OSS RPCs
 - Thousands of files aggregated with MB-sized network transfers
 - Leverage existing high throughput OSS bulk transfer rates
 - 1GB/s OSS read/write is about 30,000 32KB files/sec
- ▶ Unregister+delete container to mass-delete **all** files within
 - Simplifies user data management, accounting, job cleanup
 - Avoid MDS overhead when dealing with groups of related files

Client Container Image (CCI) (2.18+)



- ▶ Ext4 filesystem images used *ad-hoc* with Lustre today
 - **Read-only cache** of many small files manually mounted on clients
 - **Root filesystem images** for diskless clients/VMs
- ▶ **Container Image is local ext4/ldiskfs image** mounted on client
 - Directory tree (maybe millions of files) stored in one CCI Lustre file
 - **Best for self-contained workloads** (AI, Genomics, ensemble runs)
 - Can configure with job preamble script today, for read-only data

- ▶ CCI automates container image handling into Lustre TODO
 - Image is registered to Lustre directory to control future access
 - **Transparently mount** registered image **at client** on directory access
 - Image data blocks read (written) from/to OST(s) and/or client cache
 - Automatically unmount from client when idle or under contention
- ▶ Internally mount and export from MDS for multi-client access
 - Will need some modifications to CCI file to *Lustre-ize* files

Comparison and Summary of WBC vs. CCI

Metadata Writeback Cache

- Keep normal namespace
- **Transparent to users**
- Very low latency metadata operations
- **Faster single client**
- Network **batch RPCs** improves other ops
- **Lower total overhead** due to fewer layers

Client Container Image

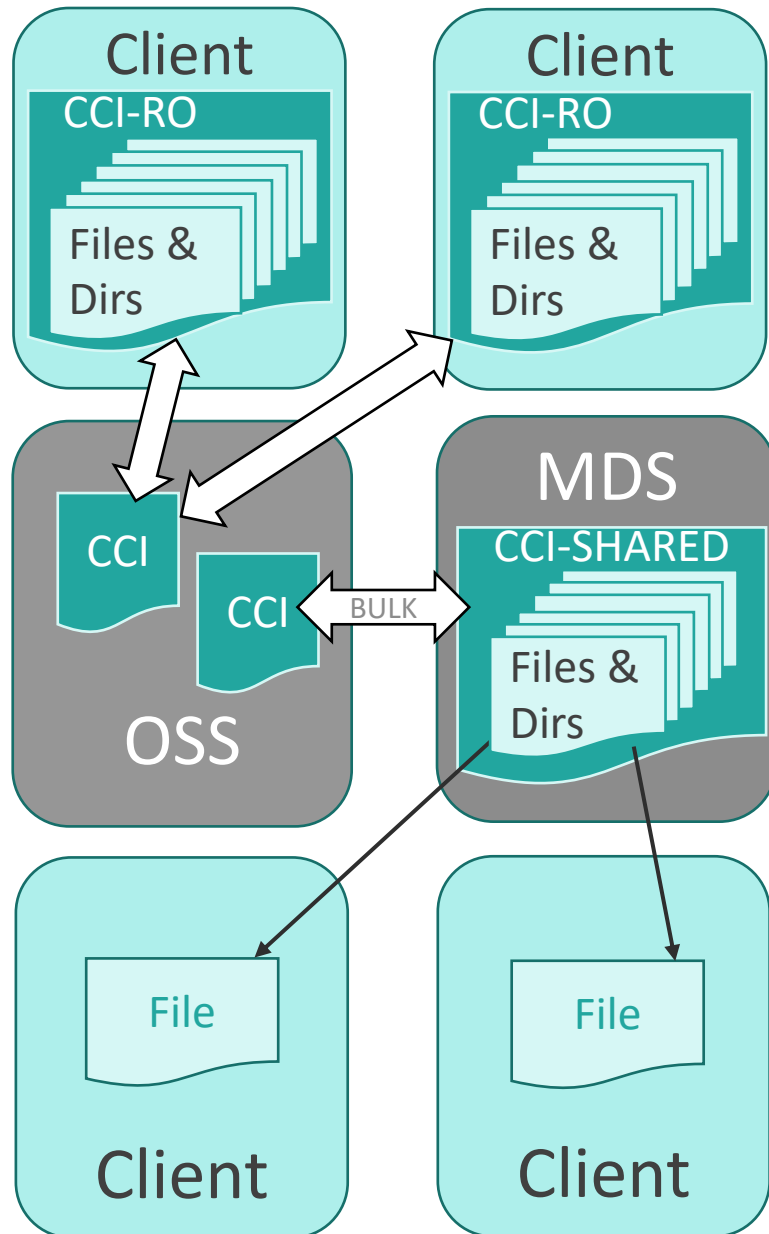
- **Segregated directory subtree**
- Needs coordination with user/job
- Not for all usage patterns
- **Faster total performance**
- Network **bulk IO** avoids MDS workload
- Bulk file/data management (e.g. fast unlink)
- **Metadata tiering/HSM aggregation**

- Significant improvements for evolving HPC workloads
- Leverages substantial functionality that already exists
- Needs supporting project to drive steady progress



Whamcloud


Thank You!
Questions?



CCI Access Models

- ▶ Need to integrate image handling on Lustre client/MDS
 - Integrate CCI creation with job workflow is easiest
 - CCI layout type on parent directory creates CCI upon mkdir
 - Enhance ldiskfs online resize to manage image size
- ▶ Client **exclusively** mounts CCI(s) and modifies locally
 - For initial image creation/import from directory tree
 - For workloads that run independently per directory
- ▶ Multiple clients **read-only** mount single image
 - Shared input datasets (e.g. gene sequence, AI training)
- ▶ MDS exports **shared read-write** image to many clients
 - Internal mount at MDS attaches image to namespace
 - Use Data-on-MDT to transparently export image tree to clients
- ▶ Process whole tree of small files for HSM/tiering
 - Efficiently migrate tree to/from flash tier, to/from archive

Single Client **32KB** File Create Performance (MDS+OSS)




1 client, n processes, 12000 files/process

`mdtest -n 12000 -d ${OUTDIR} -u -v -p 20 -w 32768 -e 32768 -F -i 3`


MDS, OSS: 6x 960GB NVMe, 2.7GHz 24-core 8186, 48GB RAM

Single Client 32KB File Create Performance (MDS vs. CCI)



► Early testing of manually-configured CCI shows significant promise

Single Client **32KB** File Create Performance (MDS vs. CCI)



- ▶ 4.15 CCI improvement due to Ubuntu 4.15 kernel loopback driver
- ▶ Early testing of CCI prototype shows promise