

内核通知链

1、背景

在linux内核中，各个子系统之间有很强的相互关系，某些子系统可能对其他子系统产生的事件比较感兴趣。因此内核引入了**notifier**机制，当然了**notifier**机制只能用在在内核子系统之间，不能用在在内核与应用层之间。比如当系统**suspend**的时候，就会使用到**notifier**机制来通知系统的内核线程进行**suspend**。内核实现的**notifier**机制代码位于**kernel/kernel/notifier.c**，同时此机制的代码量也不是很多只有**600**行左右。

2、数据结构

内核使用 `struct notifier_block` 结构代表一个 **notifier**

```
typedef int (*notifier_fn_t)(struct notifier_block *nb,
                             unsigned long action, void *data);

struct notifier_block {
    //代表事件发生之后调用的回调函数
    notifier_fn_t notifier_call;
    //用来链接同一类型的notifier
    struct notifier_block __rcu *next;
    // 通知链的优先级，数字越大优先级越高，优先执行
    int priority;
};
```

3、通知链的类型

- 原子通知链

仅仅对 `notifier_block` 的封装。`atomic_notifier_head` 中包含 `spin_lock` 表示不能睡眠通知链元素的回调函数（事件发生要执行的函数）在中断或原子操作上下文中运行，不允许阻塞

```
struct atomic_notifier_head {
    spinlock_t lock;
    struct notifier_block __rcu *head;
};
```

- 可阻塞通知链

包含读写信号量成员**rwsem**，信号量的特点就是运行在进程上下文，还可以睡眠。

```
struct blocking_notifier_head {
    struct rw_semaphore rwsem;
    struct notifier_block __rcu *head;
};
```

- 原始通知链

没有任何限制，需要调用者维护

```

struct raw_notifier_head {
    struct notifier_block __rcu *head;
};

```

- SRCU通知链

是block notifier chain的变体，采用SRCU（Sleepable Read-Copy Update）代替rw-semaphore保护chains

```

struct srcu_struct {
    short srcu_lock_nesting[2]; /* srcu_read_lock() nesting depth. */
    short srcu_idx;             /* Current reader array element. */
    u8 srcu_gp_running;         /* GP workqueue running? */
    u8 srcu_gp_waiting;         /* GP waiting for readers? */
    struct swait_queue_head srcu_wq;
                                /* Last srcu_read_unlock() wakes GP. */
    struct rcu_head *srcu_cb_head; /* Pending callbacks: Head. */
    struct rcu_head **srcu_cb_tail; /* Pending callbacks: Tail. */
    struct work_struct srcu_work; /* For driving grace periods. */
#ifdef CONFIG_DEBUG_LOCK_ALLOC
    struct lockdep_map dep_map;
#endif /* #ifdef CONFIG_DEBUG_LOCK_ALLOC */
};

struct srcu_notifier_head {
    // 通知链可能被多个线程同时访问或修改，所以需要使用互斥锁保证线程安全性
    struct mutex mutex;
    // 提供机制，在不加锁情况下读取共享数据结构，通过在更新数据时延迟释放旧版本的数据来实现
    struct srcu_struct srcu;
    struct notifier_block __rcu *head;
};

```

4、notifier chain初始化

内核提供一套宏初始化各个类型的通知链（动态初始化）

```

#define ATOMIC_INIT_NOTIFIER_HEAD(name) do { \
    spin_lock_init(&(name)->lock); \
    (name)->head = NULL; \
} while (0)
#define BLOCKING_INIT_NOTIFIER_HEAD(name) do { \
    init_rwsem(&(name)->rwsem); \
    (name)->head = NULL; \
} while (0)
#define RAW_INIT_NOTIFIER_HEAD(name) do { \
    (name)->head = NULL; \
} while (0)

```

静态初始化通知链（静态初始化）

```

#define ATOMIC_NOTIFIER_INIT(name) { \
    .lock = __SPIN_LOCK_UNLOCKED(name.lock), \
}

```

```

        .head = NULL }

#define BLOCKING_NOTIFIER_INIT(name) { \
        .rwsem = __RWSEM_INITIALIZER((name).rwsem), \
        .head = NULL }

#define RAW_NOTIFIER_INIT(name) { \
        .head = NULL }

#define SRCU_NOTIFIER_INIT(name, pcpu) \
{ \
        .mutex = __MUTEX_INITIALIZER(name.mutex), \
        .head = NULL, \
        .srcu = __SRCU_STRUCT_INIT(name.srcu, pcpu), \
}

#define ATOMIC_NOTIFIER_HEAD(name) \
    struct atomic_notifier_head name = \
        ATOMIC_NOTIFIER_INIT(name)

#define BLOCKING_NOTIFIER_HEAD(name) \
    struct blocking_notifier_head name = \
        BLOCKING_NOTIFIER_INIT(name)

#define RAW_NOTIFIER_HEAD(name) \
    struct raw_notifier_head name = \
        RAW_NOTIFIER_INIT(name)

```

SRCU 通知链不能使用静态方法，内核提供了一个动态的初始化函数

```

/**
 * srcu_init_notifier_head - Initialize an SRCU notifier head
 * @nh: Pointer to head of the srcu notifier chain
 *
 * Unlike other sorts of notifier heads, SRCU notifier heads require
 * dynamic initialization. Be sure to call this routine before
 * calling any of the other SRCU notifier routines for this head.
 *
 * If an SRCU notifier head is deallocated, it must first be cleaned
 * up by calling srcu_cleanup_notifier_head(). Otherwise the head's
 * per-cpu data (used by the SRCU mechanism) will leak.
 */
void srcu_init_notifier_head(struct srcu_notifier_head *nh)
{
    mutex_init(&nh->mutex);
    if (init_srcu_struct(&nh->srcu) < 0)
        BUG();
    nh->head = NULL;
}

```

5、注册 / 注销通知链

内核提供的最基本的注册通知链函数

通过判断 **priority** 的大小，倒序插入

```

/*
 * Notifier chain core routines. The exported routines below
 * are layered on top of these, with appropriate locking added.

```

```

*/

static int notifier_chain_register(struct notifier_block **nl,
                                  struct notifier_block *n)
{
    while ((*nl) != NULL) {
        if (unlikely((*nl) == n)) {
            WARN(1, "double register detected");
            return 0;
        }
        if (n->priority > (*nl)->priority)
            break;
        nl = &((*nl)->next);
    }
    n->next = *nl;
    rcu_assign_pointer(*nl, n);
    return 0;
}

```

注销函数：先找到此节点，然后从链表中删除一个操作。因为插入 / 删除操作都是临界资源，需要使用 **rcu** 机制保护起来。同理，内核通过包装核心的注册 / 注销函数，实现上述说的四种

notifier chain

```

extern int atomic_notifier_chain_register(struct atomic_notifier_head *nh,
                                          struct notifier_block *nb);
extern int blocking_notifier_chain_register(struct blocking_notifier_head *nh,
                                          struct notifier_block *nb);
extern int raw_notifier_chain_register(struct raw_notifier_head *nh,
                                       struct notifier_block *nb);
extern int srcu_notifier_chain_register(struct srcu_notifier_head *nh,
                                       struct notifier_block *nb);

extern int atomic_notifier_chain_unregister(struct atomic_notifier_head *nh,
                                           struct notifier_block *nb);
extern int blocking_notifier_chain_unregister(struct blocking_notifier_head
*nh,
                                           struct notifier_block *nb);
extern int raw_notifier_chain_unregister(struct raw_notifier_head *nh,
                                         struct notifier_block *nb);
extern int srcu_notifier_chain_unregister(struct srcu_notifier_head *nh,
                                         struct notifier_block *nb);

```

6、通知函数

当某种事件需要发生的时候，需要调用内核提供的通知函数 **notifier call**，来通知注册过响应的子系统

```

/**
 * notifier_call_chain - Informs the registered notifiers about an event.
 * @nl:      Pointer to head of the blocking notifier chain
 * @val:     Value passed unmodified to notifier function
 * @v:       Pointer passed unmodified to notifier function
 * @nr_to_call:  Number of notifier functions to be called. Don't care

```



```

static void notifier_test_exit(void)
{
    printk(KERN_EMERG "notifier: notifier_test_exit!\n");
    unregister_test_notifier(&test_chain_nb);
}

module_init(notifier_test_init);
module_exit(notifier_test_exit);
MODULE_LICENSE("GPL v2");

```

call.c

```

#include <linux/kernel.h>
#include <linux/module.h>
#include <linux/notifier.h>

struct blocking_notifier_head test_chain_head;

static int call_notifier_call_chain(unsigned long val)
{
    int ret = blocking_notifier_call_chain(&test_chain_head, val, NULL);
    return notifier_to_errno(ret);
}

static int call_test_init(void)
{
    printk(KERN_EMERG "notifier: call_test_init!\n");
    call_notifier_call_chain(123); //在init函数中触发事件

    return 0;
}

static void call_test_exit(void)
{
    printk(KERN_EMERG "notifier: call_test_exit!\n");
}

module_init(call_test_init);
module_exit(call_test_exit);
MODULE_LICENSE("GPL v2");

```

Makefile

```

ifneq ($(KERNELRELEASE),)
obj-m :=notifiler.o
else
KDIR :=/lib/modules/$(shell uname -r)/build
all:
    make -C $(KDIR) M=$(PWD) modules
clean:
    rm -f *.ko *.o *.mod.o *.mod.c *.symvers *.order
endif

```

